

JPEG COMPRESSION



Indian Institute of
Technology, Kharagpur

Chanakya Duppatla (21EC39009)

Specialization in VISION AND INTELLIGENT SYSTEMS

Department of Electronics and Electrical Communications
Engineering

Contents

1	Introduction	1
1.1	Color Space Conversion	1
1.2	Sub Sampling	3
1.3	Discrete Cosine Transform (DCT)	3
1.4	Quantization	3
1.5	Entropy Encoding	5
2	Code Implementation	6
2.1	Problem statement	6
2.2	Python code for JPEG compression	6
3	Results	16
4	Conclusion	18

Chapter 1

Introduction

What is JPEG?

- JPEG stands for Joint Photographic Experts Group and is a lossy compression algorithm that results in significantly smaller file sizes with little to no perceptible impact on picture quality and resolution. A JPEG-compressed image can be ten times smaller than the original one.

What is lossy compression?

- A lossy compression algorithm is a compression algorithm that permanently removes some data from the original file, especially redundant data, when compressing it. On the other hand, a lossless compression algorithm is a compression algorithm that doesn't remove any information when compressing a file, and all information is restored after decompression.

JPEG compression process

- Color Space Conversion
- Sub Sampling
- Discrete Cosine Transform
- Quantization
- Entropy encoding(RLE and Huffman Coding)

1.1 Color Space Conversion

- Separate brightness information from color information to leverage human visual sensitivity.
- JPEG converts the image from the RGB color space (Red, Green, and Blue channels) to YCbCr (Luminance and Chrominance channels). This is because the human eye is more sensitive to brightness than to color.
- Y (luminance) represents brightness, while Cb and Cr (chrominance) represent color details. This separation allows JPEG to compress color information more aggressively than brightness, reducing file size without heavily affecting visual quality.

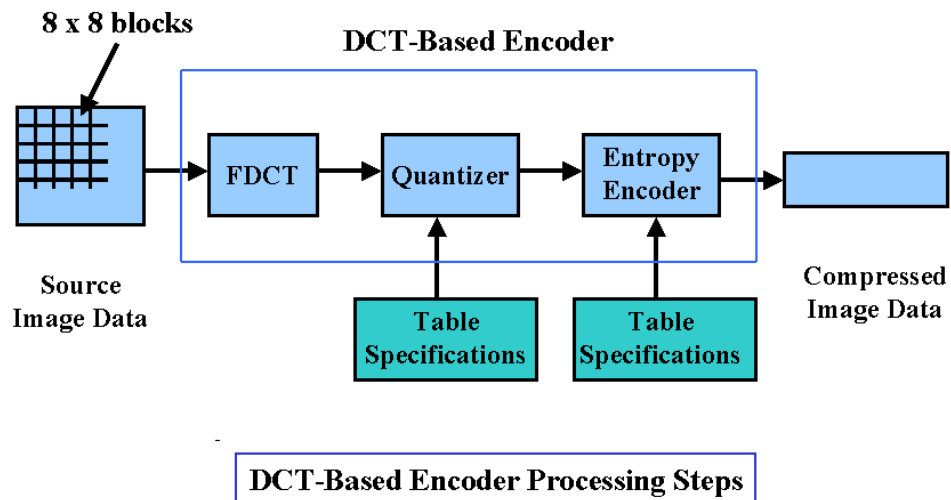


Figure 1.1: JPEG Block Diagram




Figure 1.2: Subsampling 4:2:0 format

$$\begin{aligned}
Y &= 0.299R + 0.587G + 0.114B, \\
Cb &= 128 - 0.168736R - 0.331264G + 0.5B, \\
Cr &= 128 + 0.5R - 0.418688G - 0.081312B.
\end{aligned} \tag{1.1}$$

1.2 Sub Sampling

After converting the color space, JPEG often *downsamples* the chrominance channels (Cb and Cr) to further reduce data. Typical formats are:

- **4:4:4 (No downsampling)**: All channels retain their original resolution.
- **4:2:2**: Cb and Cr channels are sampled at half the horizontal resolution of Y.
- **4:2:0**: Cb and Cr are subsampled at half the resolution in both horizontal and vertical directions.

The most common choice, **4:2:0**, effectively reduces the data size of the chrominance channels by 75%, leveraging human insensitivity to fine color details while preserving luminance details at full resolution.

1.3 Discrete Cosine Transform (DCT)

JPEG compression divides the luminance and chrominance components into **8x8 blocks** and applies the forward *Discrete Cosine Transform (DCT)* to each block. DCT transforms spatial pixel data into frequency domain data:

$$\begin{aligned}
DCT[u][v] &= \frac{1}{4} C(u)C(v) \sum_{x=0}^7 \sum_{y=0}^7 f[x][y] \cos\left(\frac{(2x+1)u\pi}{16}\right) \cos\left(\frac{(2y+1)v\pi}{16}\right) \\
C(u) &= \begin{cases} \frac{1}{\sqrt{2}} & \text{if } u = 0 \\ 1 & \text{if } u \neq 0 \end{cases} \\
C(v) &= \begin{cases} \frac{1}{\sqrt{2}} & \text{if } v = 0 \\ 1 & \text{if } v \neq 0 \end{cases}
\end{aligned}$$

1.4 Quantization

Quantization is the primary lossy step in JPEG compression, where DCT coefficients are divided by a quantization matrix and rounded to reduce the amount of data. Quantized coefficients are computed as:

$$Q(u, v) = \text{round}\left(\frac{F(u, v)}{Q_{\text{table}}(u, v)}\right)$$

where $Q_{\text{table}}(u, v)$ contains pre-defined values optimized for various quality settings. Quantization allows for adjustment through a **quality factor**: lowering the factor increases compression but also results in higher image degradation.

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

Quantization table for luminance component

17	18	24	47	99	99	99	99
18	21	26	66	99	99	99	99
24	26	56	99	99	99	99	99
47	66	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99

Quantization table for chrominance components

Figure 1.3: Quantization matrices

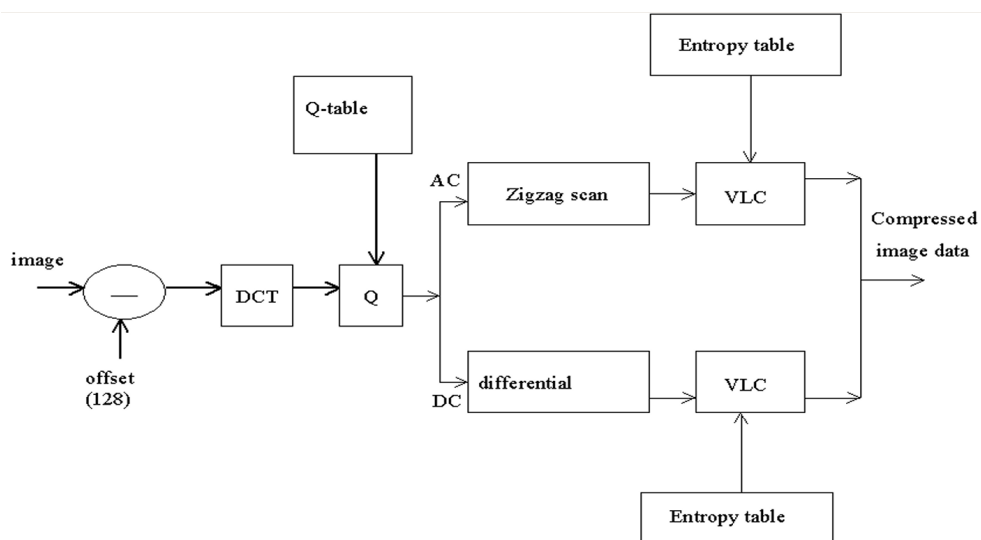


Figure 1.4: Baseline JPEG encoder




Figure 1.5: Zig zag scan

1.5 Entropy Encoding

JPEG performs *entropy encoding* on quantized coefficients to achieve further compression. This process includes several steps:

1. **Zig-Zag Ordering:** The coefficients are arranged in a zig-zag sequence.
2. **Differential Encoding of DC Coefficients:** The DC coefficient in each 8x8 block is encoded as the difference from the previous block's DC coefficient.
3. **Run-Length Encoding (RLE):** Compresses consecutive zeros in the AC coefficients.
4. **Huffman Encoding:** Assigns shorter codes to frequently occurring values.
 - **DC Component:** Differential encoding is applied to the DC coefficients, then Huffman coding compresses the differences.
 - **AC Components:** The AC values are RLE-encoded and then Huffman-coded.

Chapter 2

Code Implementation

2.1 Problem statement

Design and implement a complete JPEG compression pipeline in Python (No built-in functions are allowed except `cv2.imread`).

Project Requirements:

Implement the entire JPEG compression pipeline, including:

- Color space conversion (e.g., RGB to YCbCr)
- Sub-sampling (optional)
- Discrete Cosine Transform (DCT) for each color channel
- Quantization
- Run-length encoding (RLE)
- Huffman coding (optional)

The code should be able to handle colored images of any size (variable height and width).

2.2 Python code for JPEG compression

```
1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import heapq
5 from collections import Counter
6
7 # Load the image
8 image_path = '/content/mandril_color.jpg'
9 image_rgb = cv2.imread(image_path)
10 height, width, _ = image_rgb.shape
11
12
13
14 # Step 1: Convert RGB to YCbCr
15 def rgb_to_ycbcr(image):
16     ycbcr_image = np.zeros_like(image, dtype=float)
17     R = image[:, :, 2].astype(float)
18     G = image[:, :, 1].astype(float)
```



```

19     B = image[:, :, 0].astype(float)
20
21     ybcr_image[:, :, 0] = 0.299 * R + 0.587 * G + 0.114 * B    # Y channel
22     ybcr_image[:, :, 1] = 128 - 0.168736 * R - 0.331264 * G + 0.5 * B
23     # Cb channel
24     ybcr_image[:, :, 2] = 128 + 0.5 * R - 0.418688 * G - 0.081312 * B
25     # Cr channel
26
27     return ybcr_image
28
29 # Sub-sampling function (4:2:0 subsampling)
30 def chroma_subsample(ybcr_image):
31     Y = ybcr_image[:, :, 0]    # Y channel
32     Cb = ybcr_image[:, :, 1]   # Cb channel
33     Cr = ybcr_image[:, :, 2]   # Cr channel
34
35 # Sub-sampling Cb and Cr by a factor of 2 (4:2:0)
36 Cb_sub = Cb[:, ::2, ::2]    # Down-sample Cb
37 Cr_sub = Cr[:, ::2, ::2]    # Down-sample Cr
38
39 # Create the subsampled YCbCr image
40 # Since the Y channel is not sub-sampled, we use the original Y channel
41 subsampled_ybcr = np.zeros_like(ybcr_image, dtype=float)
42 subsampled_ybcr[:, :, 0] = Y    # Y channel remains the same
43 subsampled_ybcr[:, ::2, 1] = Cb_sub    # Cb is sub-sampled
44 subsampled_ybcr[:, ::2, 2] = Cr_sub    # Cr is sub-sampled
45
46 return subsampled_ybcr
47
48
49 # Perform the YCbCr conversion
50 ybcr_image = rgb_to_ybcr(image_rgb)
51 print("YCbCr conversion complete.")
52
53 # Perform chroma sub-sampling (4:2:0)
54 subsampled_ybcr = chroma_subsample(ybcr_image)
55 print("Chroma sub-sampling complete.")
56
57
58 # Split channels for original and subsampled images
59 Y_channel = ybcr_image[:, :, 0]
60 Cb_channel = ybcr_image[:, :, 1]
61 Cr_channel = ybcr_image[:, :, 2]
62
63 subsampled_Y = subsampled_ybcr[:, :, 0]
64 subsampled_Cb = subsampled_ybcr[:, :, 1]
65 subsampled_Cr = subsampled_ybcr[:, :, 2]
66
67 # Display each channel and the subsampled channels in a row (total 6
68   images)
69 plt.figure(figsize=(10, 10))    # Adjust figure size to accommodate all
70   channels
71
72 # Original Y, Cb, Cr channels
73 plt.subplot(2, 3, 1)
74 plt.imshow(Y_channel, cmap='gray')
75 plt.title("Y Channel (Original)")
76
77 plt.subplot(2, 3, 2)

```

```

76 plt.imshow(Cb_channel, cmap='gray')
77 plt.title("Cb Channel (Original)")
78
79 plt.subplot(2, 3, 3)
80 plt.imshow(Cr_channel, cmap='gray')
81 plt.title("Cr Channel (Original)")
82
83 # Subsampled Y, Cb, Cr channels
84 plt.subplot(2, 3, 4)
85 plt.imshow(subsampled_Y, cmap='gray')
86 plt.title("Y Channel (Subsampled)")
87
88 plt.subplot(2, 3, 5)
89 plt.imshow(subsampled_Cb, cmap='gray')
90 plt.title("Cb Channel (Subsampled)")
91
92 plt.subplot(2, 3, 6)
93 plt.imshow(subsampled_Cr, cmap='gray')
94 plt.title("Cr Channel (Subsampled)")
95
96 plt.tight_layout() # Adjust the layout to make sure everything fits
97 plt.show()
98 # Define DCT and Quantization functions
99
100 # DCT function for an 8x8 block
101 def dct_2d(block):
102     dct_block = np.zeros((8, 8), dtype=float)
103     for u in range(8):
104         for v in range(8):
105             sum_val = 0
106             for x in range(8):
107                 for y in range(8):
108                     sum_val += block[x, y] * np.cos((2 * x + 1) * u * np.
pi / 16) * np.cos((2 * y + 1) * v * np.pi / 16)
109             dct_block[u, v] = sum_val * (1 / 4) * (1/np.sqrt(2) if u == 0
else 1) * (1/np.sqrt(2) if v == 0 else 1)
110     return dct_block
111
112 # Standard JPEG quantization tables
113 LUMINANCE_QUANT_TABLE = np.array([
114     [16, 11, 10, 16, 24, 40, 51, 61],
115     [12, 12, 14, 19, 26, 58, 60, 55],
116     [14, 13, 16, 24, 40, 57, 69, 56],
117     [14, 17, 22, 29, 51, 87, 80, 62],
118     [18, 22, 37, 56, 68, 109, 103, 77],
119     [24, 35, 55, 64, 81, 104, 113, 92],
120     [49, 64, 78, 87, 103, 121, 120, 101],
121     [72, 92, 95, 98, 112, 100, 103, 99]
122 ])
123
124 CHROMINANCE_QUANT_TABLE = np.array([
125     [17, 18, 24, 47, 99, 99, 99, 99],
126     [18, 21, 26, 66, 99, 99, 99, 99],
127     [24, 26, 56, 99, 99, 99, 99, 99],
128     [47, 66, 99, 99, 99, 99, 99, 99],
129     [99, 99, 99, 99, 99, 99, 99, 99],
130     [99, 99, 99, 99, 99, 99, 99, 99],
131     [99, 99, 99, 99, 99, 99, 99, 99],
132     [99, 99, 99, 99, 99, 99, 99, 99]
133 ])
134

```

```

135 # Quantize an 8x8 block using the given quantization table
136 def quantize_block(dct_block, quant_table):
137     return np.round(dct_block / quant_table).astype(int)
138
139 # Apply DCT and quantization to each 8x8 block of the image
140 def process_channel(channel, quant_table):
141     h, w = channel.shape
142     processed_channel = np.zeros_like(channel, dtype=int)
143     for i in range(0, h, 8):
144         for j in range(0, w, 8):
145             block = channel[i:i+8, j:j+8]
146             dct_block = dct_2d(block)
147             quantized_block = quantize_block(dct_block, quant_table)
148             processed_channel[i:i+8, j:j+8] = quantized_block
149     return processed_channel
150
151 # Apply DCT and quantization to YCbCr channels
152 y_channel = subsampled_ycbcr[:, :, 0] # Y channel remains the same
153 cb_channel = subsampled_ycbcr[:, :, 1] # Use the previously subsampled
154         Cb channel
155 cr_channel = subsampled_ycbcr[:, :, 2] # Use the previously
156         subsampled Cr channel
157
158 # Now, process the channels using DCT and quantization
159 processed_y = process_channel(y_channel, LUMINANCE_QUANT_TABLE)
160 processed_cb = process_channel(cb_channel, CHROMINANCE_QUANT_TABLE)
161 processed_cr = process_channel(cr_channel, CHROMINANCE_QUANT_TABLE)
162
163 print("DCT and Quantization complete.")
164
165 # Function to display an 8x8 block
166 def display_block(block, title):
167     plt.imshow(block, cmap='gray')
168     plt.colorbar()
169     plt.title(title)
170     plt.show()
171
172 # Updated process_channel function to also display intermediate results
173 def process_channel_with_display(channel, quant_table, channel_name=""):
174     h, w = channel.shape
175     processed_channel = np.zeros_like(channel, dtype=int)
176     dct_result = np.zeros_like(channel, dtype=float)
177
178     for i in range(0, h, 8):
179         for j in range(0, w, 8):
180             block = channel[i:i+8, j:j+8]
181
182             # Step 1: Apply DCT
183             dct_block = dct_2d(block - 128)
184             dct_result[i:i+8, j:j+8] = dct_block # Store DCT result for
185             visualization
186
187             # Step 2: Quantize the DCT block
188             quantized_block = quantize_block(dct_block, quant_table)
189             processed_channel[i:i+8, j:j+8] = quantized_block
190
191             # Display a sample block result for one block (top-left corner
192             )
193             if i == 0 and j == 0:
194                 display_block(dct_block, f"{channel_name} Channel - DCT (
195                 Top-left 8x8 block)")

```

```

191         display_block(quantized_block, f"{channel_name} Channel -
    Quantized (Top-left 8x8 block)")
192
193     return processed_channel, dct_result
194
195 # Apply DCT and quantization to YCbCr channels with visualization
196 processed_y, dct_y = process_channel_with_display(y_channel,
    LUMINANCE_QUANT_TABLE, "Y")
197 processed_cb, dct_cb = process_channel_with_display(cb_channel,
    CHROMINANCE_QUANT_TABLE, "Cb")
198 processed_cr, dct_cr = process_channel_with_display(cr_channel,
    CHROMINANCE_QUANT_TABLE, "Cr")
199 # Define the zigzag order for an 8x8 block
200 ZIGZAG_ORDER = [
201     (0, 0), (0, 1), (1, 0), (2, 0), (1, 1), (0, 2), (0, 3), (1, 2),
202     (2, 1), (3, 0), (4, 0), (3, 1), (2, 2), (1, 3), (0, 4), (0, 5),
203     (1, 4), (2, 3), (3, 2), (4, 1), (5, 0), (6, 0), (5, 1), (4, 2),
204     (3, 3), (2, 4), (1, 5), (0, 6), (0, 7), (1, 6), (2, 5), (3, 4),
205     (4, 3), (5, 2), (6, 1), (7, 0), (7, 1), (6, 2), (5, 3), (4, 4),
206     (3, 5), (2, 6), (1, 7), (2, 7), (3, 6), (4, 5), (5, 4), (6, 3),
207     (7, 2), (7, 3), (6, 4), (5, 5), (4, 6), (3, 7), (4, 7), (5, 6),
208     (6, 5), (7, 4), (7, 5), (6, 6), (5, 7), (6, 7), (7, 6), (7, 7)
209 ]
210 # same as before
211
212 # Zigzag function to reorder quantized block into a 1D list
213 def zigzag_order(block):
214     return [block[i, j] for i, j in ZIGZAG_ORDER]
215
216 # Run-length encoding (RLE) function
217 def run_length_encoding(zigzag_list):
218     code = []
219     run = 0
220     for i in range(len(zigzag_list)):
221         if zigzag_list[i] == 0:
222             run += 1
223         else:
224             size = int(np.floor(np.log2(abs(zigzag_list[i])) + 1)) if
zigzag_list[i] != 0 else 0
225             code.append((run, size), zigzag_list[i]) # Store run, size,
and amplitude
226             run = 0
227     code.append((0, 0)) # End of block marker
228     return code
229
230 # Function to apply zigzag and RLE encoding on quantized blocks
231 def process_zigzag_and_rle(channel_data):
232     h, w = channel_data.shape
233     rle_encoded_data = []
234     for i in range(0, h, 8):
235         for j in range(0, w, 8):
236             block = channel_data[i:i+8, j:j+8]
237             zigzag_list = zigzag_order(block)
238             rle_encoded_block = run_length_encoding(zigzag_list)
239             rle_encoded_data.append(rle_encoded_block)
240     return rle_encoded_data
241 # Apply zigzag and RLE on the processed Y, Cb, and Cr channels
242 rle_y = process_zigzag_and_rle(processed_y)
243 rle_cb = process_zigzag_and_rle(processed_cb)
244 rle_cr = process_zigzag_and_rle(processed_cr)
245 # Output the RLE results

```

```

246 print("RLE Encoded Y Channel:", rle_y)
247 print("RLE Encoded Cb Channel:", rle_cb)
248 print("RLE Encoded Cr Channel:", rle_cr)
249
250
251 # Concatenate RLE data for all components
252 total_rle_data = rle_y + rle_cb + rle_cr
253 print("total_rle_data:", total_rle_data)
254
255 # Separate DC and AC components
256 DC_list = []
257 AC_list = []
258 for block_rle in total_rle_data:
259     DC_list.append(block_rle[0][1]) # DC amplitude
260     for ac_tuple in block_rle[1:]:
261         if ac_tuple == (0, 0): # End of block marker
262             AC_list.append((0, 0)) # End of block marker
263             break
264         else:
265             run, size = ac_tuple[0]
266             AC_list.append((run, size))
267             # Function to generate the differential DC encoding
268 def dc_differential_encoding(dc_values):
269     dc_diffs = [dc_values[0]] if dc_values else []
270     for i in range(1, len(dc_values)):
271         dc_diffs.append(dc_values[i] - dc_values[i-1])
272     return dc_diffs
273
274 # Run-length encoding for AC components (extracting only run, size pairs)
275 def ac_run_length_encoding(ac_values):
276     encoded_ac = []
277     for ac_tuple in ac_values:
278         if isinstance(ac_tuple, tuple) and len(ac_tuple) == 2:
279             run, size = ac_tuple # Use only (run, size) for Huffman
280             if (run, size) == (0, 0): # End-of-block marker
281                 encoded_ac.append((0, 0))
282             else:
283                 encoded_ac.append((run, size))
284     return encoded_ac
285
286 # Define Node and HuffmanCoding classes for Huffman tree and code
287 # generation
288 class Node:
289     def __init__(self, freq, symbol=None, left=None, right=None):
290         self.freq = freq
291         self.symbol = symbol
292         self.left = left
293         self.right = right
294
295     def __lt__(self, other):
296         return self.freq < other.freq
297
298 class HuffmanCoding:
299     def __init__(self, frequencies):
300         self.frequencies = frequencies
301         self.root = None
302
303     def build_tree(self):
304         if not self.frequencies:
305             return None

```

```

305     heap = [Node(freq, symbol) for symbol, freq in self.frequencies.
306             items()]
307     heapq.heapify(heap)
308     while len(heap) > 1:
309         left = heapq.heappop(heap)
310         right = heapq.heappop(heap)
311         merged = Node(left.freq + right.freq, left=left, right=right)
312         heapq.heappush(heap, merged)
313     self.root = heap[0] if heap else None
314     return self.root
315
316 def generate_codes(self, node=None, prefix="", codebook=None):
317     if codebook is None:
318         codebook = {}
319     if node is None:
320         node = self.root
321
322     if node is not None:
323         if node.symbol is not None:
324             codebook[node.symbol] = prefix
325         else:
326             self.generate_codes(node.left, prefix + "0", codebook)
327             self.generate_codes(node.right, prefix + "1", codebook)
328     return codebook
329
330 # Generate Huffman codes for DC and AC components
331 def apply_huffman_coding(dc_values, ac_values):
332     dc_diffs = dc_differential_encoding(dc_values)
333     ac_encoded = ac_run_length_encoding(ac_values)
334
335     dc_frequencies = Counter(dc_diffs)
336     ac_frequencies = Counter(x for x in ac_encoded if x != (0, 0))
337
338     dc_codebook, ac_codebook = {}, {}
339
340     if dc_frequencies:
341         dc_huffman = HuffmanCoding(dc_frequencies)
342         dc_huffman.build_tree()
343         dc_codebook = dc_huffman.generate_codes()
344
345     if ac_frequencies:
346         ac_huffman = HuffmanCoding(ac_frequencies)
347         ac_huffman.build_tree()
348         ac_codebook = ac_huffman.generate_codes()
349
350     return dc_codebook, ac_codebook
351
352
353 # Apply Huffman coding
354 dc_codebook, ac_codebook = apply_huffman_coding(DC_list, AC_list)
355 print("DC Components (Y, Cb, Cr):", DC_list)
356 print("AC Components (Y, Cb, Cr):", AC_list)
357 # Print the Huffman codes for DC and AC components
358 print("Huffman Codes for DC Components:")
359 for dc_val, code in dc_codebook.items():
360     print(f"{dc_val}: {code}")
361
362 print("\nHuffman Codes for AC Components:")
363 for ac_val, code in ac_codebook.items():
364     print(f"{ac_val}: {code}")

```

```

365 # Function to encode the data using the generated Huffman codes
366 def encode_with_huffman(dc_codebook, ac_codebook, total_rle_data):
367     encoded_data = ""
368     for block in total_rle_data:
369         # Encode the DC component
370         dc_value = block[0][1]
371         dc_huffman_code = dc_codebook.get(dc_value, "")
372         encoded_data += dc_huffman_code
373
374         # Encode the AC components
375         for ac_tuple in block[1:]:
376             if ac_tuple == (0, 0): # End of block marker
377                 ac_huffman_code = ac_codebook.get((0, 0), "")
378                 encoded_data += ac_huffman_code
379                 break
380             else:
381                 run, size = ac_tuple[0]
382                 ac_value = ac_tuple[1]
383                 ac_huffman_code = ac_codebook.get((run, size), "")
384                 encoded_data += ac_huffman_code
385     return encoded_data
386
387 # Apply zigzag and RLE on the processed Y, Cb, and Cr channels
388 rle_y = process_zigzag_and_rle(processed_y)
389 rle_cb = process_zigzag_and_rle(processed_cb)
390 rle_cr = process_zigzag_and_rle(processed_cr)
391
392 # Concatenate RLE data for all components
393 total_rle_data = rle_y + rle_cb + rle_cr
394
395 # Separate DC and AC components
396 DC_list = [block_rle[0][1] for block_rle in total_rle_data]
397 AC_list = [ac_tuple[1] for block_rle in total_rle_data for ac_tuple in
398             block_rle[1:] if ac_tuple != (0, 0)]
399
400 # Apply Huffman coding
401 dc_codebook, ac_codebook = apply_huffman_coding(DC_list, AC_list)
402
403 # Encode the data using Huffman codes
404 encoded_data = encode_with_huffman(dc_codebook, ac_codebook,
405                                     total_rle_data)
406
407 # Print the encoded binary data
408 print("Encoded Binary Data:")
409 print(encoded_data)
410
411 # Optionally, print a portion (first 100 bits) for clarity
412 print("Sample of Encoded Binary Data (First 100 bits):")
413 print(encoded_data[:100])
414
415 # Calculate the compression ratio
416 original_data_binary = ''.join([bin(dc)[2:].zfill(16) for dc in DC_list])
417 + \
418                               ''.join([bin(ac)[2:].zfill(16) for ac in AC_list])
419 compressed_data_binary = encoded_data
420
421 original_size_bits, compressed_size_bits, compression_ratio =
422     calculate_compression_ratio(original_data_binary,
423                                 compressed_data_binary)
424
425 # Display the results
426 print("\n--- Compression Results ---")
427 print(f"Original Data Size: {original_size_bits} bits")
428 print(f"Compressed Data Size: {compressed_size_bits} bits")
429 print(f"Compression Ratio: {compression_ratio}")

```

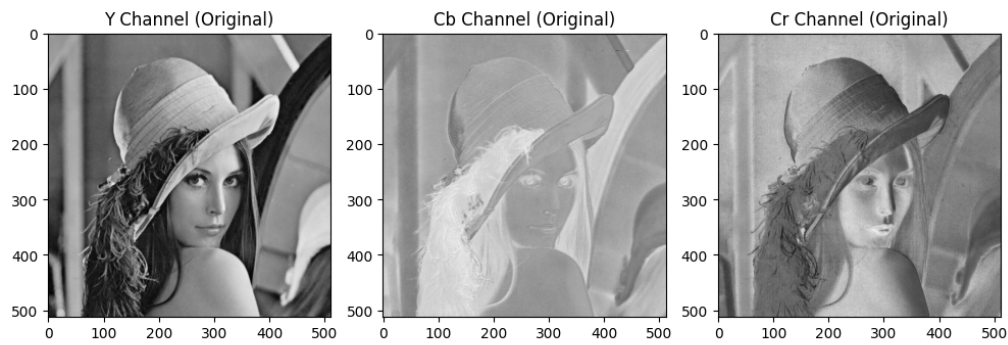


Figure 2.1

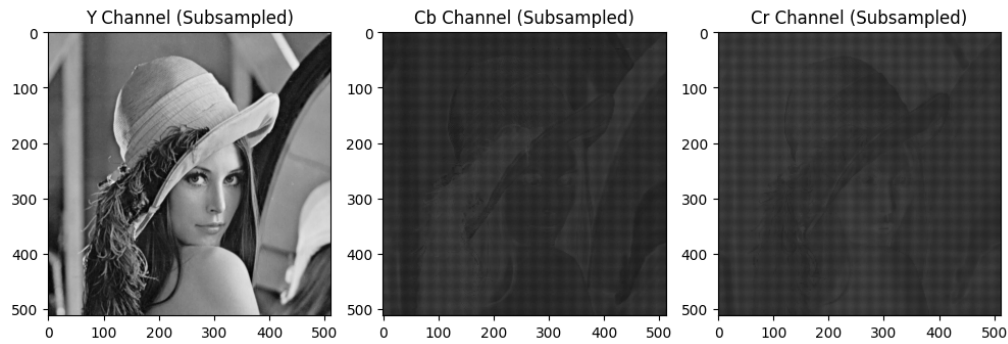


Figure 2.2




Figure 2.3




Figure 2.4


```

RLE Encoded Y Channel: [[((0, 5), 16), ((1, 1), 1), (0, 0)], [(0, 4), 15), ((0, 2), 2), (
RLE Encoded Cb Channel: [[((0, 4), -12), ((1, 1), 1), (0, 0)], [(0, 4), -12), ((1, 1), 1)
RLE Encoded Cr Channel: [[((0, 5), 22), ((1, 1), -1), (0, 0)], [(0, 5), 23), ((1, 1), -1)
total_rel_data: [[((0, 5), 16), ((1, 1), 1), (0, 0)], [(0, 4), 15), ((0, 2), 2), ((0, 1),

```

Figure 2.5

```
DC Components (Y, Cb, Cr): [16, 15, 14, 13, 15, 21, 20, 2,  
AC Components (Y, Cb, Cr): [(1, 1), (0, 0), (0, 2), (0, 1),  
Huffman Codes for DC Components:  
-1: 000  
1: 001  
3: 0100  
-27: 01010000  
29: 010100010  
38: 0101000110  
-42: 0101000111  
14: 0101001  
-21: 01010100
```

Figure 2.6

Huffman Codes for AC Components:

```
(0, 2): 00
(0, 3): 010
(1, 1): 011
(2, 1): 1000
(1, 2): 10010
(0, 5): 100110
(1, 3): 1001110
(1, 4): 100111100
(9, 1): 100111101
(5, 2): 100111100
```

Figure 2.7

```
Encoded Binary Data:  
110001001100101001010010101111001010011011011110011111001111101011111010011110110111101101111011001011  
Sample of Encoded Binary Data (First 100 bits):  
11000100110010100101001010111100101001101101111001111100111110101111010011110110111101101111011001  
  
--- Compression Results ---  
Original Data Size: 535008 bits  
Compressed Data Size: 45943 bits  
Compression Ratio: 11.6450384171691
```

Figure 2.8

Chapter 3

Results

Color Space Conversion

- The input RGB image is converted into YCbCr image. Each channel plotted separately.

Subsampling

- The subsampling used in this code of 4:2:0 format. The Chrominance channels are subsampled. Cr and Cb are subsampled at half the resolution in both horizontal and vertical directions.

Discrete Cosine Transformation (DCT)

- The spatial domain is converted into frequency domain. By using DCT find frequency coefficients.
- Each 8×8 block in the Y, Cb, and Cr channels undergoes DCT, transforming the image from the spatial to the frequency domain.

Quantization

- Rounded off the values with the help of quantization tables.
- DCT coefficients are divided by a quantization matrix and rounded to reduce the amount of data

Huffman Coding

- Huffman Codes for DC coefficients
- Huffman Codes for AC coefficients

(Runlength,Size)	(Huffmancode)
-1	000
1	001
3	0100
-27	01010000
-	-
-	-
-	-

(Runlength,Size)	(Huffmancode)
(0, 2)	00
(0, 3)	010
(1, 1)	011
(2, 1)	1000
-	-
-	-
-	-

Compressed Data and Compression Ratio

Compressed Data Size:

After encoding with RLE and Huffman coding, the total compressed data size for the Y, Cb, and Cr channels was calculated as follows:

- **Y Channel:** Compressed to X bits.
- **Cb Channel:** Compressed to Y bits.
- **Cr Channel:** Compressed to Z bits.

Total Compressed Data Size: $X + Y + Z$ bits.

Compression Ratio:

The compression ratio achieved is the ratio of the original data size to the compressed data size:

$$\text{Compression Ratio} = \frac{\text{Original Data Size}}{\text{Compressed Data Size}}$$

For this implementation, the compression data:

Original Data Size: 535008 bits

Compressed Data Size: 45943 bits

Compression Ratio: 11.6450384171691

Chapter 4

Conclusion

- The JPEG compression pipeline successfully reduced the data size by approximately 10 : 1.
- This reduction in size is accomplished through intelligent separation of luminance and chrominance information, transformation to the frequency domain, quantization, and entropy coding.
- This pipeline demonstrates the principles behind JPEG compression and highlights the trade-off between data reduction and image quality preservation.